

# Churn problem

## Import necessary libraries

```
In [1]: import numpy as np
```

```
In [2]: import pandas as pd
```

```
In [3]: import seaborn as sns
```

```
In [4]: import matplotlib.ticker as mtick
```

```
In [5]: import matplotlib.pyplot as plt
```

```
In [6]: sns.set(style = 'white')
```

```
In [7]: import os
```

```
In [8]: from sklearn.preprocessing import StandardScaler
```

```
In [9]: from sklearn.linear_model import LogisticRegression
```

```
In [10]: from sklearn import metrics
```

```
In [11]: from sklearn.ensemble import RandomForestClassifier
```

```
In [12]: from sklearn.svm import SVC
```

```
In [13]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [14]: from sklearn.tree import DecisionTreeClassifier
```

```
In [15]: from xgboost import XGBClassifier
```

```
In [16]: from sklearn.metrics import classification_report, confusion_matrix

In [17]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict

In [18]: from sklearn.preprocessing import KBinsDiscretizer
from imblearn.pipeline import make_pipeline
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn.ensemble import (RandomForestClassifier, AdaBoostClassifier,
GradientBoostingClassifier, ExtraTreesClassifier)

In [19]: from sklearn.metrics import accuracy_score, roc_curve, f1_score, precision_score, recall_score, confusion_matrix
from sklearn.metrics import roc_auc_score

In [22]: #Load the dataset
telecom_cust = pd.read_csv("C:\\Users\\30698\\Downloads\\churn-train.csv")

In [23]: telecom_cust.head()
```

Out[23]:

vice	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
DSL	Yes	No	Yes	No	No	No	Month-to-month	No	'Bank transfer (automatic)'	33.60	2117.2	No
iptic'	No	Yes	Yes	Yes	No	No	'Two year'	No	'Bank transfer (automatic)'	90.45	6565.85	No
iptic'	No	No	No	No	Yes	No	Month-to-month	Yes	'Electronic check'	84.00	424.75	No
DSL	Yes	Yes	Yes	Yes	No	No	'Two year'	No	'Bank transfer (automatic)'	67.40	3306.85	No
No	'No internet service'	'No internet service'	'No internet service'	'No internet service'	'No internet service'	'No internet service'	Month-to-month	Yes	'Bank transfer (automatic)'	19.70	168.9	No

Data Preprocessing

```
In [24]: telecom_cust.columns.values
```

```
Out[24]: array(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',  
              'PhoneService', 'MultipleLines', 'InternetService',  
              'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',  
              'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',  
              'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',  
              'TotalCharges', 'Churn'], dtype=object)
```

```
In [25]: telecom_cust.dtypes
```

```
Out[25]: gender          object  
SeniorCitizen      int64  
Partner            object  
Dependents          object  
tenure              int64  
PhoneService        object  
MultipleLines        object  
InternetService      object  
OnlineSecurity        object  
OnlineBackup          object  
DeviceProtection      object  
TechSupport           object  
StreamingTV           object  
StreamingMovies        object  
Contract              object  
PaperlessBilling      object  
PaymentMethod          object  
MonthlyCharges        float64  
TotalCharges          object  
Churn                  object  
dtype: object
```

```
In [27]: #convert Total charges to numeric variable
telecom_cust.TotalCharges = pd.to_numeric(telecom_cust.TotalCharges, errors='coerce')
print(telecom_cust.isnull().values.any())
telecom_cust.isnull().sum()
#there are 6 missing values in the dataset
```

True

```
Out[27]: gender          0
SeniorCitizen          0
Partner                0
Dependents             0
tenure                 0
PhoneService           0
MultipleLines          0
InternetService        0
OnlineSecurity         0
OnlineBackup           0
DeviceProtection       0
TechSupport            0
StreamingTV            0
StreamingMovies        0
Contract               0
PaperlessBilling       0
PaymentMethod          0
MonthlyCharges         0
TotalCharges           6
Churn                  0
dtype: int64
```

```
In [29]: #for test set
test = pd.read_csv("C:\\Users\\30698\\Downloads\\churn-test.csv")
test.TotalCharges = pd.to_numeric(test.TotalCharges, errors='coerce')
print(test.isnull().values.any())
test.isnull().sum()
```

True

```
Out[29]: gender          0
SeniorCitizen          0
Partner                0
Dependents              0
tenure                 0
PhoneService           0
MultipleLines           0
InternetService        0
OnlineSecurity         0
OnlineBackup           0
DeviceProtection       0
TechSupport            0
StreamingTV            0
StreamingMovies        0
Contract               0
PaperlessBilling        0
PaymentMethod          0
MonthlyCharges         0
TotalCharges           5
Churn                  0
dtype: int64
```

```
In [30]: #removing missing values
telecom_cust.dropna(inplace = True)
test.dropna(inplace = True)
```

```
In [31]: print(telecom_cust.isnull().values.any())
```

False

```
In [32]: df=telecom_cust
```

```
In [33]: df['Churn'].replace(to_replace='Yes', value=1, inplace=True)
df['Churn'].replace(to_replace='No', value=0, inplace=True)
```

```
In [34]: test['Churn'].replace(to_replace='Yes', value=1, inplace=True)
test['Churn'].replace(to_replace='No', value=0, inplace=True)
```

```
In [37]: #convert categorical variables to dummies
df_dummies = pd.get_dummies(df)
df_dummies.head()
```

Out[37]:


	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	Partner_Yes	Dependents_No	...	StreamingMovies_Yes	Contract_'One year'	Contract_'Two year'
0	0	61	33.60	2117.20	0	0	1	0	1	0	...	0	0	0
1	0	72	90.45	6565.85	0	0	1	0	1	0	...	0	0	0
2	0	5	84.00	424.75	0	1	0	1	0	1	...	0	0	0
3	0	49	67.40	3306.85	0	1	0	1	0	1	...	0	0	0
4	0	8	19.70	168.90	0	0	1	1	0	1	...	0	0	0

5 rows × 46 columns

```
In [38]: test_dummies = pd.get_dummies(test)
```

```
In [39]: plt.figure(figsize=(15,8))
df_dummies.corr()['Churn'].sort_values(ascending = False).plot(kind='bar')
```

Out[39]: <AxesSubplot:>



We observe a positive correlation of Churn with month-to-month contract, no tech support, no online security and fiber optic internet service. This means that most churn customers are those who have contract only for a month (this is logical as after the end of their contract is most possible to leave their "product" than others who have longer contracts), and those who don't have services like tech support and online security and also those who use fiber optic internet service because probably it is more expensive. We also observe that churn customers are negatively correlated with tenure and two year contract. This means that a customer who uses the "product" for a long time or has a two year contract is less possible to become a churn customer.

```
In [35]: #Data Analysis
```

```
In [41]: df['gender'].value_counts()*100.0 / len(df)
#Data are balanced regarding the gender (about half male, half female)
```

```
Out[41]: Male      50.245255
Female    49.754745
Name: gender, dtype: float64
```

```
In [42]: df['SeniorCitizen'].value_counts()*100.0 / len(telecom_cust)
#Most of the customers in the dataset are younger people.
```

```
Out[42]: 0      83.663894
1      16.336106
Name: SeniorCitizen, dtype: float64
```



```
In [43]: df['Dependents'].value_counts()*100.0 /len(telecom_cust)
#only 30% of customers have dependents
```

```
Out[43]: No      70.804009
         Yes      29.195991
         Name: Dependents, dtype: float64
```

```
In [44]: df['Partner'].value_counts()*100.0 /len(telecom_cust)
#half of customers have partner
```

```
Out[44]: No      52.121988
         Yes      47.878012
         Name: Partner, dtype: float64
```


```
In [45]: df['Churn'].value_counts()*100.0 /len(telecom_cust)
#The dataset is imbalanced.Only 26% are churn customers.
```

```
Out[45]: 0      73.299211
         1      26.700789
         Name: Churn, dtype: float64
```

```
In [46]: import warnings
         warnings.filterwarnings('ignore')
```


```
In [47]: ax = sns.distplot(df['tenure'], hist=True, kde=False,
                        bins=int(180/5), color = 'darkblue',
                        hist_kws={'edgecolor':'black'},
                        kde_kws={'linewidth': 4})
ax.set_ylabel('# of Customers')
ax.set_xlabel('Tenure (months)')
ax.set_title('# of Customers by their tenure')
#A lot of customers stayed only one month with the company and many stayed for 72 months.
```

Out[47]: Text(0.5, 1.0, '# of Customers by their tenure')



```
In [48]: ax = df['Contract'].value_counts().plot(kind = 'bar',rot = 0, width = 0.3)
ax.set_ylabel('# of Customers')
ax.set_title('# of Customers by Contract Type')
#Most of the customers are in the one month contract
```

Out[48]: Text(0.5, 1.0, '# of Customers by Contract Type')



```
In [49]: #The tenure of customers based on their contract type
fig, (ax1,ax2,ax3) = plt.subplots(nrows=1, ncols=3, sharey = True, figsize = (20,6))


ax = sns.distplot(df[df['Contract']=='Month-to-month']['tenure'],
                  hist=True, kde=False,
                  bins=int(180/5), color = 'turquoise',
                  hist_kws={'edgecolor':'black'},
                  kde_kws={'linewidth': 4},
                  ax=ax1)
ax.set_ylabel('# of Customers')
ax.set_xlabel('Tenure (months)')
ax.set_title('Month to Month Contract')

ax = sns.distplot(df[df['Contract']=="'One year'"]['tenure'],
                  hist=True, kde=False,
                  bins=int(180/5), color = 'steelblue',
                  hist_kws={'edgecolor':'black'},
                  kde_kws={'linewidth': 4},
                  ax=ax2)
ax.set_xlabel('Tenure (months)',size = 14)
ax.set_title('One Year Contract',size = 14)

ax = sns.distplot(df[df['Contract']=="'Two year'"]['tenure'],
                  hist=True, kde=False,
                  bins=int(180/5), color = 'darkblue',
                  hist_kws={'edgecolor':'black'},
                  kde_kws={'linewidth': 4},
                  ax=ax3)

ax.set_xlabel('Tenure (months)')
ax.set_title('Two Year Contract')
#We observe that most of the customers with the one month contract stay just for one month in the telecom company.On the other
#hand most of those with the two year contract stay for 72 months in the company.
#So,customer taking a longer contract are more loyal to the company and tend to stay with it for a longer period of time.
```

Out[49]: Text(0.5, 1.0, 'Two Year Contract')




```
In [50]: df['Contract'].unique()
```

```
Out[50]: array(['Month-to-month', "'Two year'", "'One year'"], dtype=object)
```

```
In [51]: #Churn vs Tenure
sns.boxplot(x = df.Churn, y = df.tenure)
#Loyal customers tend to stay for a longer tenure with the company.
```

```
Out[51]: <AxesSubplot:xlabel='Churn', ylabel='tenure'>
```



```

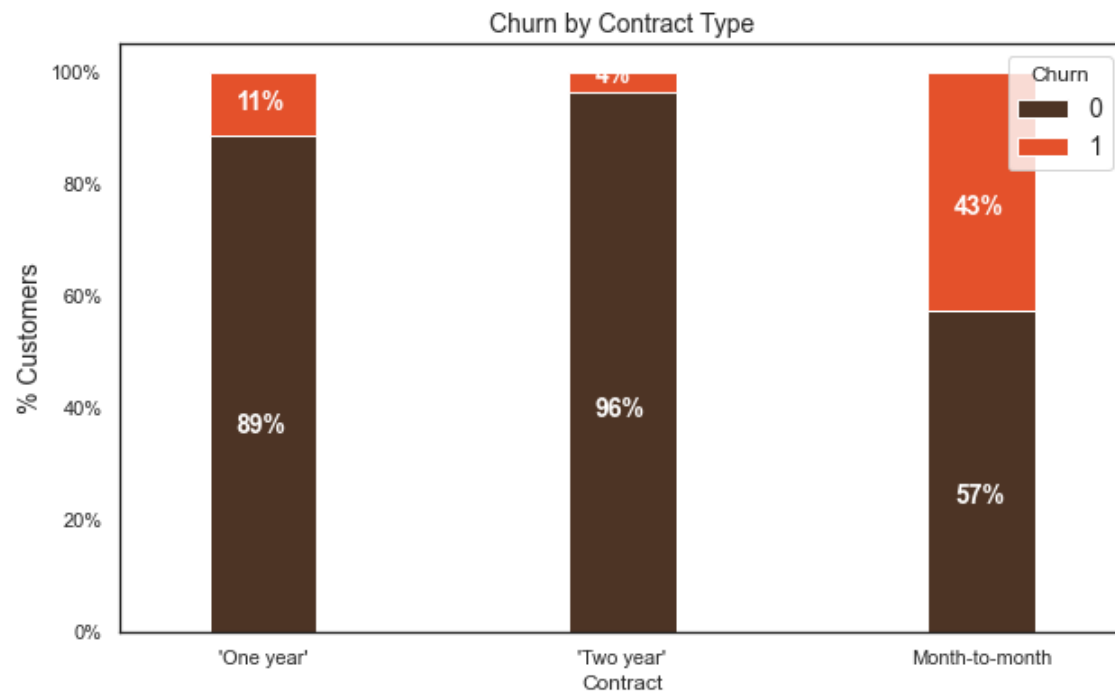
In [52]: colors = ['#4D3425', '#E4512B']
contract_churn = df.groupby(['Contract', 'Churn']).size().unstack()

ax = (contract_churn.T*100.0 / contract_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.3,
                                                                stacked = True,
                                                                rot = 0,
                                                                figsize = (10,6),
                                                                color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='best',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers',size = 14)
ax.set_title('Churn by Contract Type',size = 14)

# Code to add the data labels on the stacked bar chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.4*height),
                color = 'white',
                weight = 'bold',
                size = 14)
#Customers who have the one month contract have a very high churn rate.

```



```


In [53]: colors = ['#4D3425', '#E4512B']
seniority_churn = df.groupby(['SeniorCitizen', 'Churn']).size().unstack()

ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.2,
                                                                stacked = True,
                                                                rot = 0,
                                                                figsize = (8,6),
                                                                color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Seniority Level',size = 14)

# Code to add the data labels on the stacked bar chart
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.4*height),
                color = 'white',
                weight = 'bold',size = 14)
#Senior Citizens have higher churn rate than younger customers.(15%more)

```



```

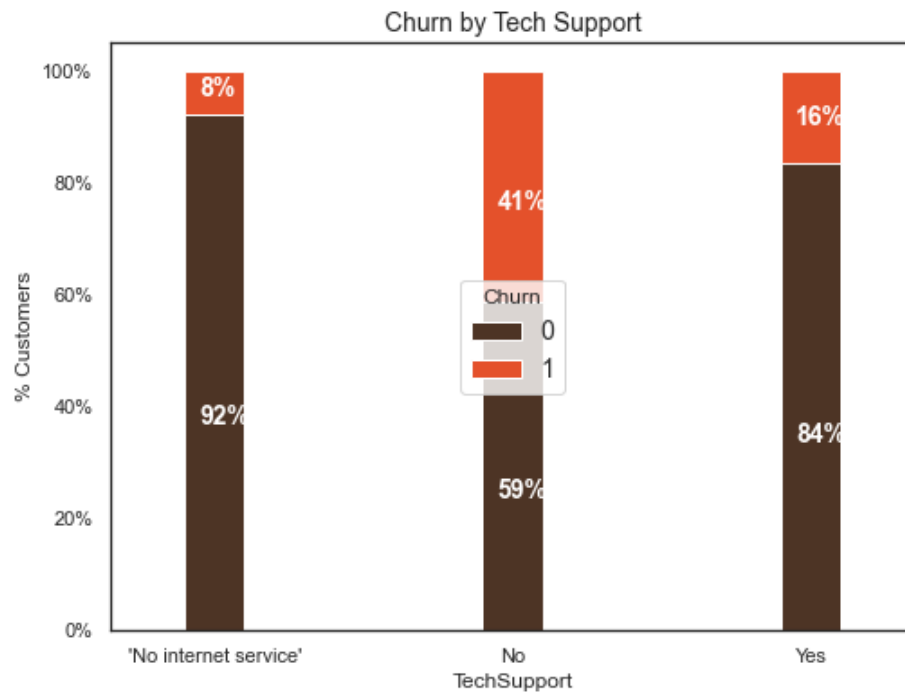
In [54]: colors = ['#4D3425', '#E4512B']
seniority_churn = df.groupby(['TechSupport', 'Churn']).size().unstack()

ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.2,
                                                                stacked = True,
                                                                rot = 0,
                                                                figsize = (8,6),
                                                                color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Tech Support',size = 14)

for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.4*height),
                color = 'white',
                weight = 'bold',size =14)
#The absence of Tech support increases significantly the churn rate.

```





```

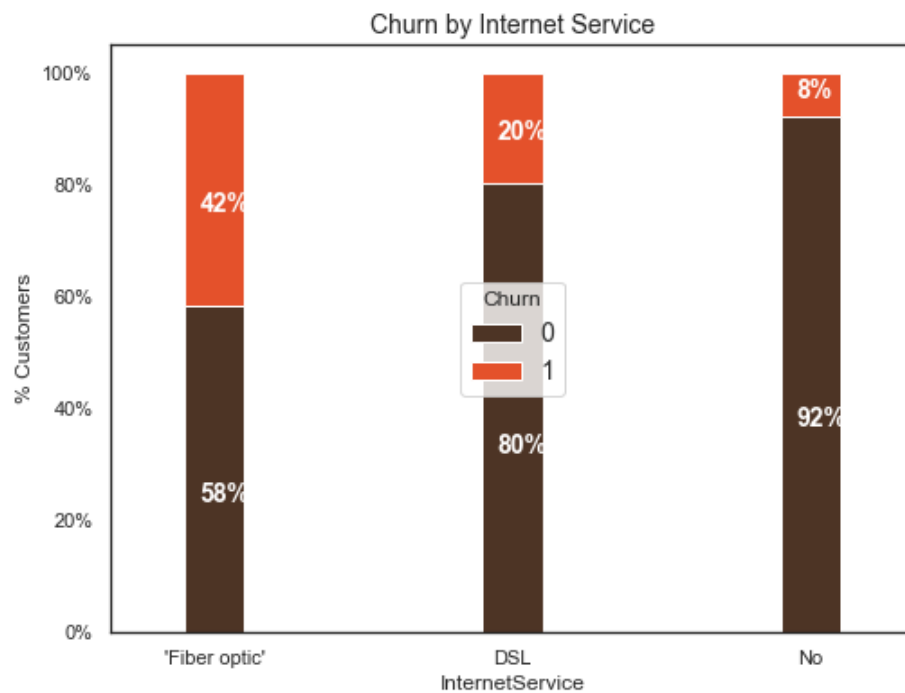
In [55]: colors = ['#4D3425', '#E4512B']
seniority_churn = df.groupby(['InternetService', 'Churn']).size().unstack()

ax = (seniority_churn.T*100.0 / seniority_churn.T.sum()).T.plot(kind='bar',
                                                                width = 0.2,
                                                                stacked = True,
                                                                rot = 0,
                                                                figsize = (8,6),
                                                                color = colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='center',prop={'size':14},title = 'Churn')
ax.set_ylabel('% Customers')
ax.set_title('Churn by Internet Service',size = 14)

for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x()+.25*width, p.get_y()+.4*height),
                color = 'white',
                weight = 'bold',size = 14)
#The use of the fiber optic service increases significantly the churn rate.

```




```
In [56]: df_dummies.head()
```

Out[56]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	Partner_Yes	Dependents_No	...	StreamingMovies_Yes	Contract_'One year'	Contract_'Two year'
0	0	61	33.60	2117.20	0	0	1	0	1	0	...	0	0	0
1	0	72	90.45	6565.85	0	0	1	0	1	0	...	0	0	0
2	0	5	84.00	424.75	0	1	0	1	0	1	...	0	0	0
3	0	49	67.40	3306.85	0	1	0	1	0	1	...	0	0	0
4	0	8	19.70	168.90	0	0	1	1	0	1	...	0	0	0

5 rows × 46 columns



```
In [57]: #There is high range in the values so we need to scale them in order to avoid the high values to prevail in our model.  
scal=StandardScaler()
```

```
In [58]: col=['tenure', 'MonthlyCharges', 'TotalCharges']
```

```
In [59]: df_dummies[col]=scal.fit_transform(df_dummies[col])
```

```
In [60]: test_dummies[col]=scal.fit_transform(test_dummies[col])
```

```
In [63]: y=df_dummies['Churn']  
X=df_dummies.drop('Churn',axis=1)
```

```
In [65]: #sampling.  
#We will use the oversampling technique because the dataset is class imbalanced.  
from imblearn.over_sampling import SMOTE  
smote = SMOTE()  
x_smote, y_smote = smote.fit_sample(X,y)
```

```
In [68]: ada = AdaBoostClassifier()
ada.fit(X,y)
scores = cross_val_score(ada, X, y, cv=10, scoring='accuracy')
#We could change the scoring to recall but we will take it separately later
print(scores)
print(scores.mean())
pred=cross_val_predict(ada,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
#Using Recall for imbalanced data
recall_score(y,pred)
```

```
[0.82089552 0.79317697 0.79744136 0.79530917 0.79530917 0.79317697
 0.79957356 0.78464819 0.81449893 0.80128205]
0.7995311902028319
[[3087  350]
 [ 590  662]]
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	3437
1	0.65	0.53	0.58	1252
accuracy			0.80	4689
macro avg	0.75	0.71	0.73	4689
weighted avg	0.79	0.80	0.79	4689

```
Out[68]: 0.5287539936102237
```

In [135]: *#The performance of adaboost classifier using the data after oversampling increases dramatically the recall score.*

```
ada2 = AdaBoostClassifier()
ada2.fit(x_smote, y_smote)
scores = cross_val_score(ada2, x_smote, y_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(ada2,x_smote, y_smote, cv=10)
print(confusion_matrix(y_smote,pred))
print(classification_report(y_smote, pred))
recall_score(y_smote,pred)
```

```
[0.7747093  0.7630814  0.73837209 0.76744186 0.81659389 0.82387191
 0.84425036 0.82678311 0.84861718 0.83697234]
```

```
0.8040693443011409
```

```
[[2655  782]
```

```
 [ 565 2872]]
```

	precision	recall	f1-score	support
0	0.82	0.77	0.80	3437
1	0.79	0.84	0.81	3437
accuracy			0.80	6874
macro avg	0.81	0.80	0.80	6874
weighted avg	0.81	0.80	0.80	6874

Out[135]: 0.8356124527203957


## Feature selection

```
In [69]: rf = RandomForestClassifier(n_estimators=100)
rf.fit(X,y)
rf_features=rf.feature_importances_
sorted_idx = np.argsort(rf_features)
print(sorted_idx )
plt.figure(figsize=(15, 15))
plt.barh(range(len(sorted_idx)), rf_features[sorted_idx], align='center')
plt.yticks(range(len(sorted_idx)), X.columns.values[sorted_idx])
plt.xlabel('Importance')
plt.title('Feature importances')
plt.draw()
plt.show()

#Here we can see the most important and less important features.Random Forest does feature selection
#so we dont have to do cross validation during the training of this model.
#As we can see the most important features to predict a churn customer are:Total charges,tenure,monthly charges,
#month to month contract,absence of tech support,fiber optic internet service.
#we can drop the Least important features:'PhoneService_No', 'PhoneService_Yes', 'MultipleLines_'No phone service'',
# 'InternetService_No', 'OnlineSecurity_'No internet service'', 'OnlineBackup_'No internet service'',
# 'DeviceProtection_'No internet service'' 'TechSupport_'No internet service'', 'StreamingTV_'No internet service'',
# 'StreamingMovies_'No internet service''
```


```
[21 18 17 30 24 10 33 11 12 27 29 31 16 36 26 35 34 41 44  9 23 32 20 42
  8 14 13 25  6  7 39 40 22  5 37  4  0 28 15 19 43 38  2  1  3]
```

Feature importances



```
In [70]: fs_dataset=df_dummies.drop(columns=['PhoneService_No','PhoneService_Yes', "MultipleLines_'No phone service'", 'InternetService_No', "OnlineSecurity_StreamingMovies_'No internet service'"])
y_f=fs_dataset['Churn']
X_f=fs_dataset.drop('Churn',axis=1)
```

```
In [131]: from sklearn.feature_selection import RFE
rfe = RFE(estimator=DecisionTreeClassifier(), n_features_to_select=15)
rfe.fit(X,y)
rfe.fit(X_f,y_f)
```



```
In [82]: from sklearn.ensemble import ExtraTreesClassifier
from sklearn.feature_selection import SelectFromModel
clf = ExtraTreesClassifier(n_estimators=15)
clf = clf.fit(X, y)
model = SelectFromModel(clf, prefit=True)
```

```
In [87]: pip=make_pipeline(model,smote,ada)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
recall_score(y,y_predict)
```

0.7197301795082549

```
[[2524 239]
 [ 913 1013]]
```

	precision	recall	f1-score	support
0	0.73	0.91	0.81	2763
1	0.81	0.53	0.64	1926
accuracy			0.75	4689
macro avg	0.77	0.72	0.73	4689
weighted avg	0.77	0.75	0.74	4689

Out[87]: 0.8091054313099042

```
In [80]: pip=make_pipeline(rfe,smote,ada)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
recall_score(y,y_predict)
```

```
0.7261116364326875
```

```
[[2641  283]
```

```
 [ 796  969]]
```

	precision	recall	f1-score	support
0	0.77	0.90	0.83	2924
1	0.77	0.55	0.64	1765
accuracy			0.77	4689
macro avg	0.77	0.73	0.74	4689
weighted avg	0.77	0.77	0.76	4689

```
Out[80]: 0.7739616613418531
```



```

In [81]: gr = GradientBoostingClassifier()
gr.fit(X,y)
scores = cross_val_score(gr, X,y, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(gr,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
recall_score(y,pred)

[0.79957356 0.79957356 0.7782516  0.78464819 0.81663113 0.78678038
 0.80383795 0.78464819 0.80383795 0.81623932]
0.7974021832230788
[[3103  334]
 [ 615  637]]

```

	precision	recall	f1-score	support
0	0.83	0.90	0.87	3437
1	0.66	0.51	0.57	1252
accuracy			0.80	4689
macro avg	0.75	0.71	0.72	4689
weighted avg	0.79	0.80	0.79	4689

Out[81]: 0.5087859424920128

```
In [76]: #after sampling
gr = GradientBoostingClassifier()
gr.fit(x_smote,y_smote)
scores = cross_val_score(gr, x_smote,y_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(gr,x_smote,y_smote, cv=10)
print(confusion_matrix(y_smote,pred))
print(classification_report(y_smote, pred))
recall_score(y_smote,pred)
```

```
[0.78052326 0.77180233 0.74709302 0.77034884 0.8588064 0.86899563
 0.88646288 0.86899563 0.88646288 0.86608443]
0.8305575302122474
[[2775 662]
 [ 503 2934]]
```

	precision	recall	f1-score	support
0	0.85	0.81	0.83	3437
1	0.82	0.85	0.83	3437
accuracy			0.83	6874
macro avg	0.83	0.83	0.83	6874
weighted avg	0.83	0.83	0.83	6874

Out[76]: 0.853651440209485

```
In [99]: pip=make_pipeline(rfe,smote,gr)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
recall_score(y,y_predict)
```

```
0.7354042156547894
[[2689  280]
 [ 748  972]]
      precision    recall  f1-score   support

      0       0.78      0.91      0.84      2969
      1       0.78      0.57      0.65      1720

   accuracy          0.78          0.78      4689
  macro avg       0.78      0.74      0.75      4689
weighted avg       0.78      0.78      0.77      4689
```

Out[99]: 0.7763578274760383

```
In [88]: pip=make_pipeline(model,smote,gr)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
recall_score(y,y_predict)
```

```
0.7386242291455964
[[2663  252]
 [ 774 1000]]
      precision    recall  f1-score   support

      0       0.77      0.91      0.84      2915
      1       0.80      0.56      0.66      1774

   accuracy          0.78          0.78      4689
  macro avg       0.79      0.74      0.75      4689
weighted avg       0.78      0.78      0.77      4689
```

Out[88]: 0.7987220447284346

```
In [125]: rf=RandomForestClassifier()
rf.fit(X,y)
scores = cross_val_score(rf, X,y, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(rf,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
recall_score(y,pred)
```

```
[0.7771855  0.76226013 0.7761194  0.80063966 0.80576307]
```

```
0.7843935528941662
```

```
[[3082  355]
```

```
 [ 652  600]]
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	3437
1	0.63	0.48	0.54	1252
accuracy			0.79	4689
macro avg	0.73	0.69	0.70	4689
weighted avg	0.77	0.79	0.78	4689

```
Out[125]: 0.4792332268370607
```

```
In [126]: rf=RandomForestClassifier()
rf.fit(x_smote,y_smote)
scores = cross_val_score(rf, x_smote,y_smote, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(rf,x_smote,y_smote, cv=10)
print(confusion_matrix(y_smote,pred))
print(classification_report(y_smote, pred))
recall_score(y_smote,pred)
```

```
[0.78690909 0.784      0.88872727 0.90472727 0.91048035]
```

```
0.8549687971417228
```

```
[[2892  545]
```

```
 [ 420 3017]]
```

	precision	recall	f1-score	support
0	0.87	0.84	0.86	3437
1	0.85	0.88	0.86	3437
accuracy			0.86	6874
macro avg	0.86	0.86	0.86	6874
weighted avg	0.86	0.86	0.86	6874

```
Out[126]: 0.8778004073319755
```

```
In [93]: clf = LogisticRegression(class_weight="balanced")
clf.fit(X,y)
scores = cross_val_score(clf, X,y, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
#we also need to use the right metric for imbalanced data(recall)
recall_score(y,pred)
```

```
[0.75692964 0.74840085 0.74626866 0.71002132 0.72921109 0.73987207
 0.76119403 0.75053305 0.76972281 0.74358974]
0.7455743261713411
[[2504  933]
 [ 260  992]]
```

	precision	recall	f1-score	support
0	0.91	0.73	0.81	3437
1	0.52	0.79	0.62	1252
accuracy			0.75	4689
macro avg	0.71	0.76	0.72	4689
weighted avg	0.80	0.75	0.76	4689

Out[93]: 0.792332268370607

```
In [94]: pip=make_pipeline(rfe,smote,clf)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
```

```
0.71129221295567
[[2516  264]
 [ 921  988]]
```

	precision	recall	f1-score	support
0	0.73	0.91	0.81	2780
1	0.79	0.52	0.63	1909
accuracy			0.75	4689
macro avg	0.76	0.71	0.72	4689
weighted avg	0.76	0.75	0.73	4689

```
0.7891373801916933
```

```
In [95]: pip=make_pipeline(model,smote,clf)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
```

0.7079926836844121

[[2519 277]

[ 918 975]]

	precision	recall	f1-score	support
0	0.73	0.90	0.81	2796
1	0.78	0.52	0.62	1893
accuracy			0.75	4689
macro avg	0.76	0.71	0.71	4689
weighted avg	0.75	0.75	0.73	4689

0.7787539936102237

```
In [102]: #knn suffers from imbalanced data (knn is a distance based model).We will use weighted knn.#The intuition behind
#weighted kNN, is to give more weight to the points which are nearby and less weight to the points
#which are farther away.So the performance of the model improves when k increases using weighted knn.
```

```
In [96]: k=[1,5,11,21]
for i in k:
    neigh=KNeighborsClassifier(n_neighbors=i,weights='distance')
    neigh.fit(X,y)
    scores = cross_val_score(neigh, X, y, cv=10, scoring='accuracy')
    print(scores)
    print(scores.mean())
    pred=cross_val_predict(neigh, X,y, cv=10)
    print(confusion_matrix(y,pred))
    print(classification_report(y, pred))
    print(recall_score(y,pred))
```

```
[0.69296375 0.70362473 0.70149254 0.69722814 0.68230277 0.68230277
 0.73134328 0.7249467 0.73987207 0.75854701]
```

```
0.7114623767608843
```

```
[[2711 726]
```

```
[ 627 625]]
```

	precision	recall	f1-score	support
0	0.81	0.79	0.80	3437
1	0.46	0.50	0.48	1252
accuracy			0.71	4689
macro avg	0.64	0.64	0.64	4689
weighted avg	0.72	0.71	0.71	4689

```
0.4992012779552716
```

```
[0.75266525 0.75266525 0.73134328 0.74200426 0.75906183 0.7249467
 0.7761194 0.74413646 0.77185501 0.77777778]
```

```
0.7532575219142383
```

```
[[2892 545]
```

```
[ 612 612]]
```



```
In [98]: clf = DecisionTreeClassifier(criterion='entropy')
clf.fit(X,y)
scores = cross_val_score(clf, X, y, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
#Decision trees is not good for Large data.Decission Trees tend to overfit when there is a Large number of splits
recall_score(y,pred)
```

```
[0.74840085 0.73987207 0.74200426 0.70149254 0.71002132 0.72921109
 0.75053305 0.73560768 0.73347548 0.73504274]
0.7325661071929729
[[2757  680]
 [ 594  658]]
      precision    recall  f1-score   support

      0         0.82        0.80        0.81        3437
      1         0.49        0.53        0.51        1252

 accuracy          0.73        4689
 macro avg         0.66        0.66        0.66        4689
 weighted avg         0.73        0.73        0.73        4689
```

Out[98]: 0.5255591054313099

```
In [106]: from sklearn import svm
clf = svm.SVC(kernel='linear')
clf.fit(X,y)
scores = cross_val_score(clf, X,y, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
recall_score(y,pred)
```

```
[0.82302772 0.78678038 0.76972281 0.79317697 0.80810235 0.7803838
 0.80170576 0.79957356 0.78678038 0.8034188 ]
0.7952672534762086
[[3072  365]
 [ 595  657]]
```

	precision	recall	f1-score	support
0	0.84	0.89	0.86	3437
1	0.64	0.52	0.58	1252
accuracy			0.80	4689
macro avg	0.74	0.71	0.72	4689
weighted avg	0.79	0.80	0.79	4689

Out[106]: 0.5247603833865815

```
In [108]: from sklearn import svm
clf2 = svm.SVC(kernel='linear')
clf2.fit(x_smote,y_smote)
scores = cross_val_score(clf2,x_smote,y_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf2,x_smote,y_smote, cv=10)
print(confusion_matrix(y_smote,pred))
print(classification_report(y_smote, pred))
recall_score(y_smote,pred)
```

```
[0.72819767 0.73546512 0.7122093  0.7877907  0.90247453 0.90684134
 0.91703057 0.90684134 0.92139738 0.90684134]
```

```
0.842508928269185
```

```
[[2971  466]
```

```
 [ 617 2820]]
```

	precision	recall	f1-score	support
0	0.83	0.86	0.85	3437
1	0.86	0.82	0.84	3437
accuracy			0.84	6874
macro avg	0.84	0.84	0.84	6874
weighted avg	0.84	0.84	0.84	6874

```
Out[108]: 0.8204829793424498
```

```
In [107]: pip=make_pipeline(rfe,smote,clf)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
accuracy_score(y,y_predict)
```

0.6960097263387475

[[2248 193]

[1189 1059]]

	precision	recall	f1-score	support
0	0.65	0.92	0.76	2441
1	0.85	0.47	0.61	2248
accuracy			0.71	4689
macro avg	0.75	0.70	0.69	4689
weighted avg	0.75	0.71	0.69	4689

0.8458466453674122

Out[107]: 0.7052676476860738

```
In [109]: pip=make_pipeline(model,smote,clf)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
accuracy_score(y,y_predict)
```

0.6946770521073926

[[2211 183]

[1226 1069]]

	precision	recall	f1-score	support
0	0.64	0.92	0.76	2394
1	0.85	0.47	0.60	2295
accuracy			0.70	4689
macro avg	0.75	0.69	0.68	4689
weighted avg	0.75	0.70	0.68	4689

0.8538338658146964

Out[109]: 0.6995094902964385

```
In [110]: from sklearn.naive_bayes import GaussianNB
clf = GaussianNB()
clf.fit(X,y)
scores = cross_val_score(clf, X,y, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf,X,y, cv=10)
print(confusion_matrix(y,pred))
print(classification_report(y, pred))
recall_score(y,pred)
```

[0.71855011 0.73347548 0.6652452 0.66098081 0.66950959 0.68230277  
0.71002132 0.67803838 0.70788913 0.6965812 ]  
0.6922593989758169  
[[2205 1232]  
[ 211 1041]]

	precision	recall	f1-score	support
0	0.91	0.64	0.75	3437
1	0.46	0.83	0.59	1252
accuracy			0.69	4689
macro avg	0.69	0.74	0.67	4689
weighted avg	0.79	0.69	0.71	4689

Out[110]: 0.8314696485623003

```
In [105]: clf2 = GaussianNB()
clf2.fit(x_smote,y_smote)
scores = cross_val_score(clf2, x_smote,y_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf2,x_smote,y_smote, cv=10)
print(confusion_matrix(y_smote,pred))
print(classification_report(y_smote, pred))
recall_score(y_smote,pred)
```

```
[0.77616279 0.74854651 0.72093023 0.73255814 0.75691412 0.75691412
 0.77292576 0.73508006 0.77438137 0.76128093]
0.7535694035408417
```

```
[[2269 1168]
 [ 526 2911]]
```

	precision	recall	f1-score	support
0	0.81	0.66	0.73	3437
1	0.71	0.85	0.77	3437
accuracy			0.75	6874
macro avg	0.76	0.75	0.75	6874
weighted avg	0.76	0.75	0.75	6874

```
Out[105]: 0.8469595577538551
```

```
In [111]: pip=make_pipeline(rfe,smote,clf)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
accuracy_score(y,y_predict)
```

0.7070689302071212

[[2550 297]

[ 887 955]]

	precision	recall	f1-score	support
0	0.74	0.90	0.81	2847
1	0.76	0.52	0.62	1842
accuracy			0.75	4689
macro avg	0.75	0.71	0.71	4689
weighted avg	0.75	0.75	0.74	4689

0.762779552715655

Out[111]: 0.7474941352100661

```
In [ ]: #Recall is the metric that we have to take into account in this specific problem..
#we are interested in the proportion of churns identified correctly by the total number of churns
```



```
In [116]: clf = LogisticRegression(class_weight="balanced")
clf.fit(X_f,y_f)
scores = cross_val_score(clf,X_f,y_f, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf,X_f,y_f, cv=10)
print(confusion_matrix(y_f,pred))
print(classification_report(y_f, pred))
recall_score(y_f,pred)
#(logistic regression with data from which we dropped the least important
#features using random forest "feature selection")
```

```
[0.75692964 0.74626866 0.74626866 0.70788913 0.72921109 0.73773987
 0.76119403 0.75053305 0.76972281 0.74358974]
0.7449346673227271
[[2501  936]
 [ 260  992]]
```

	precision	recall	f1-score	support
0	0.91	0.73	0.81	3437
1	0.51	0.79	0.62	1252
accuracy			0.74	4689
macro avg	0.71	0.76	0.72	4689
weighted avg	0.80	0.74	0.76	4689

```
Out[116]: 0.792332268370607
```

```
In [117]: pip=make_pipeline(rfe,smote,clf)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
#good model
```

```
0.711134676684494
[[2530  272]
 [ 907  980]]

      precision    recall  f1-score   support

     0       0.74      0.90      0.81      2802
     1       0.78      0.52      0.62      1887

 accuracy          0.75      4689
 macro avg       0.76      0.71      0.72      4689
weighted avg       0.75      0.75      0.74      4689

0.7827476038338658
```

```
In [119]: gr = GradientBoostingClassifier()
gr.fit(X_f,y_f)
scores = cross_val_score(gr, X_f,y_f, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(clf,X_f,y_f, cv=10)
print(confusion_matrix(y_f,pred))
print(classification_report(y_f, pred))
recall_score(y_f,pred)
```

```
[0.80170576 0.80597015 0.78464819 0.78464819 0.81876333 0.79530917
 0.80383795 0.78464819 0.80383795 0.81623932]
0.799960818617535
[[2501  936]
 [ 260  992]]

      precision    recall  f1-score   support

     0       0.91      0.73      0.81      3437
     1       0.51      0.79      0.62      1252

 accuracy          0.74      4689
 macro avg       0.71      0.76      0.72      4689
weighted avg       0.80      0.74      0.76      4689
```

Out[119]: 0.792332268370607

```
In [121]: pip=make_pipeline(model,smote,gr)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
```

0.7400622670052407

[[2660 245]

[ 777 1007]]

	precision	recall	f1-score	support
0	0.77	0.92	0.84	2905
1	0.80	0.56	0.66	1784
accuracy			0.78	4689
macro avg	0.79	0.74	0.75	4689
weighted avg	0.79	0.78	0.77	4689

0.8043130990415336

```
In [122]: pip=make_pipeline(rfe,smote,gr)
pip.fit(X,y)
y_predict = pip.predict(X)
accuracy_score(pip.predict(X), y)
print(roc_auc_score(pip.predict(X),y))
print(confusion_matrix(pip.predict(X), y))
print(classification_report(pip.predict(X), y))
print(recall_score(y,y_predict))
```

0.7450380143889006

[[2798 321]

[ 639 931]]

	precision	recall	f1-score	support
0	0.81	0.90	0.85	3119
1	0.74	0.59	0.66	1570
accuracy			0.80	4689
macro avg	0.78	0.75	0.76	4689
weighted avg	0.79	0.80	0.79	4689

0.7436102236421726

```
In [132]: #Now,we will use the data set after smote sampling and after the drop of the Least important
#features that was found with the Random Forest Algorithm
xf_smote, yf_smote = smote.fit_sample(X_f,y_f)
```

```
In [133]: gr = GradientBoostingClassifier()
gr.fit(xf_smote,yf_smote)
scores = cross_val_score(gr, xf_smote,yf_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(gr,xf_smote,yf_smote, cv=10)
print(confusion_matrix(yf_smote,pred))
print(classification_report(yf_smote, pred))
recall_score(yf_smote,pred)
```

```
[0.78633721 0.75872093 0.74854651 0.76744186 0.86608443 0.8558952
 0.88500728 0.86899563 0.86608443 0.86754003]
```

```
0.8270653498527469
```

```
[[2740  697]
```

```
 [ 492 2945]]
```

	precision	recall	f1-score	support
0	0.85	0.80	0.82	3437
1	0.81	0.86	0.83	3437
accuracy			0.83	6874
macro avg	0.83	0.83	0.83	6874
weighted avg	0.83	0.83	0.83	6874

```
Out[133]: 0.8568519057317427
```

```
In [134]: ada = AdaBoostClassifier()
ada.fit(xf_smote, yf_smote)
scores = cross_val_score(ada, xf_smote, yf_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(ada,xf_smote, y_smote, cv=10)
print(confusion_matrix(yf_smote,pred))
print(classification_report(yf_smote, pred))
recall_score(yf_smote,pred)
```

```
[0.79069767 0.7630814 0.74709302 0.74709302 0.82532751 0.82387191
 0.85298399 0.83697234 0.83988355 0.83551674]
```

```
0.8062521157035982
```

```
[[2643 794]
```

```
[ 538 2899]]
```

	precision	recall	f1-score	support
0	0.83	0.77	0.80	3437
1	0.78	0.84	0.81	3437
accuracy			0.81	6874
macro avg	0.81	0.81	0.81	6874
weighted avg	0.81	0.81	0.81	6874

Out[134]: 0.8434681408204829

```
In [127]: # using the following model in the test set
BEST=make_pipeline(model,smote,gr)
BEST.fit(X,y)
y_predict = BEST.predict(X)
accuracy_score(BEST.predict(X), y)
print(roc_auc_score(BEST.predict(X),y))
print(confusion_matrix(BEST.predict(X), y))
print(classification_report(BEST.predict(X), y))
print(recall_score(y,y_predict))
```

```
0.7368794188261255
```

```
[[2648 249]
```

```
[ 789 1003]]
```

	precision	recall	f1-score	support
0	0.77	0.91	0.84	2897
1	0.80	0.56	0.66	1792
accuracy			0.78	4689
macro avg	0.79	0.74	0.75	4689
weighted avg	0.78	0.78	0.77	4689

```
0.8011182108626198
```

```
In [136]: BEST2 = GradientBoostingClassifier()
BEST2.fit(xf_smote,yf_smote)
scores = cross_val_score(BEST2, xf_smote,yf_smote, cv=10, scoring='accuracy')
print(scores)
print(scores.mean())
pred=cross_val_predict(BEST2,xf_smote,yf_smote, cv=10)
print(confusion_matrix(yf_smote,pred))
print(classification_report(yf_smote, pred))
recall_score(yf_smote,pred)
```

```
[0.78633721 0.75872093 0.74854651 0.76744186 0.86608443 0.8558952
 0.88500728 0.86899563 0.86608443 0.86899563]
```

```
0.8272109102603162
```

```
[[2740  697]
```

```
 [ 492 2945]]
```

	precision	recall	f1-score	support
0	0.85	0.80	0.82	3437
1	0.81	0.86	0.83	3437
accuracy			0.83	6874
macro avg	0.83	0.83	0.83	6874
weighted avg	0.83	0.83	0.83	6874

```
Out[136]: 0.8568519057317427
```

```
In [128]: #selecting 300 customers randomly
test_data_300 = test_dummies.sample(n=300)
churn_predictions = BEST.predict(test_data_300.drop('Churn', axis=1))
number_true_churn_customers = np.count_nonzero(churn_predictions == test_data_300['Churn'])
print('The number of true churn customers among the selected 300 customers is:', number_true_churn_customers)
```

```
The number of true churn customers among the selected 300 customers is: 209
```

```
In [129]: print(sum(churn_predictions))
```

```
142
```

```
In [ ]: #Expected(cost)=10*142+64*(209-142)=5708 EUROS
```

```
In [138]: #SECOND WAY(with BEST2)
test2=test_dummies.drop(columns=['PhoneService_No','PhoneService_Yes', "MultipleLines_'No phone service'", 'InternetService_No', "OnlineSecurity_'No internet service'"])
```

```
In [139]: #selecting 300 customers randomly
test2_data_300 = test2.sample(n=300)
churn_predictions = BEST2.predict(test2_data_300.drop('Churn', axis=1))
number_true_churn_customers = np.count_nonzero(churn_predictions == test2_data_300['Churn'])
print('The number of true churn customers among the selected 300 customers is:', number_true_churn_customers)
```

The number of true churn customers among the selected 300 customers is: 232

```
In [140]: print(sum(churn_predictions))
#9282Euros the cost with the second way
```

103

## Results

**Descriptive Task** Characteristics of loyal and churn customers Customers taking a longer contract(2 years contract)are more loyal to the company and tend to stay with it for a longer period of time.Customers who dont churn tend to stay for a longer tenure with the telecom company Customers who have month to month contract have a very high churn rate. Senior citizens have almost double the churn rate than younger population When monthly charges are high a big percent of customers churn.The absence of online security is also a reason that make customers to churn. Most churn customers are those who have contract only for a month (this is logical as after the end of their contract is most possible to leave their "product" than others who have longer contracts), and those who dont have services like tech support and online security and also those who use fiber optic internet servise(which surprisingly means faster internet) because probably it is more expensive and also because the most of churn customers are senior citizens who probably dont need very fast internet. Proposed system The system I would suggest to the Telecom company is to always make a two-year contract with its customers, to have DSL internet service and to include tech support in the contract.

**Predictive Task** The predictive model used for our test set is Gradient Boosting Algorithm with smote sampling and tree based feature selection. The expected cost for the company in one month will be 5709 Euros